



Table of Contents

Lab 02: Gradient Descent

1. Loss landscape
2. The “Gradient” in Gradient Descent

3. Forward & Backward

- 3.1. Forward
- 3.2. Backward
4. Implementation
 - 4.1. Import library
 - 4.2. Create data
 - 4.3. Training

Lab 02: Gradient Descent

Copyright © Department of Computer Science, University of Science, Vietnam National University, Ho Chi Minh City

- Student name: Hoang Anh Tra
- ID: 21127453

How to do your homework

- You will work directly on this notebook; the word **TODO** indicates the parts you need to do.
- You can discuss the ideas as well as refer to the documents, but *the code and work must be yours*.

How to submit your homework

- Before submitting, save this file as `<ID>.jl`. For example, if your ID is 123456, then your file will be `123456.jl`. And export to PDF with name `123456.pdf` then submit zipped source code and pdf into `123456.zip` onto Moodle.

Note

Note that you will get 0 point for the wrong submit.

Content of the assignment:

- Gradient Descent

1. Loss landscape




Figure 1. Loss landscape visualized as a 2D plot. Source: codecamp.vn

The gradient descent method is an iterative optimization algorithm that operates over a loss landscape (also called an optimization surface). As we can see, our loss landscape has many peaks and valleys based on which values our parameters take on. Each peak is a local maximum that represents very high regions of loss – the local maximum with the largest loss across the entire loss landscape is the global maximum. Similarly, we also have local minimum which represents many small regions of loss. The local minimum with the smallest loss across the loss landscape is our global minimum. In an ideal world, we would like to find this global minimum, ensuring our parameters take on the most optimal possible values.

Each position along the surface corresponds to a particular loss value given a set of parameters \mathbf{W} (weight matrix) and \mathbf{b} (bias vector). Our goal is to try different values of \mathbf{W} and \mathbf{b} , evaluate their loss, and then take a step towards more optimal values that (ideally) have lower loss.

2. The “Gradient” in Gradient Descent

We can use \mathbf{W} and \mathbf{b} and to compute a loss function L or we are able to find our relative position on the loss landscape, but **which direction** we should take a step to move closer to the minimum.

- All We need to do is follow the slope of the gradient $\nabla_{\mathbf{W}}$. We can compute the gradient $\nabla_{\mathbf{W}}$ across all dimensions using the following equation:

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

- But, this equation has 2 problems:
 - 1. It's an **approximation** to the gradient.
 - 2. It's painfully slow.

In practice, we use the **analytic gradient** instead.

3. Forward & Backward

In this section, you will be asked to fill in the blank to form the forward process and backward process with the data defined as follows:

- Feature: \mathbf{X} (shape: $n \times d$, be already used bias trick)
- Label: \mathbf{y} (shape: $n \times 1$)
- Weight: \mathbf{W} (shape: $d \times 1$)

3.1. Forward

TODO: Consider one sample \mathbf{x}_i . Fill in the blank

$$h_i = \mathbf{x}_i^T \mathbf{W} \Rightarrow \frac{\partial h_i}{\partial \mathbf{W}} = \mathbf{x}_i$$

$$\hat{y}_i = \sigma(h_i) \Rightarrow \frac{\partial \hat{y}_i}{\partial h_i} = \sigma(h_i)(1 - \sigma(h_i))$$

$$loss_i = (\hat{y}_i - y_i)^2 \Rightarrow \frac{\partial loss_i}{\partial \hat{y}_i} = 2(\hat{y}_i - y_i)$$

3.2. Backward

Our loss function is MSE:

$$Loss = \frac{1}{n} \sum_{i=1}^n loss_i = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

Goal: Compute $\nabla Loss = \frac{\partial Loss(W)}{\partial W}$

How to compute $\nabla Loss$?: Use Chain-rule. Your work is to fill in the blank

TODO: Fill in the blank

$$\begin{aligned} \nabla Loss &= \frac{\partial Loss(W)}{\partial W} = \frac{1}{n} \sum_{i=1}^n \frac{\partial loss_i(W)}{\partial W} \\ &= \frac{1}{n} \sum_{i=1}^n 2(\hat{y}_i(W) - y_i) * \frac{\partial \hat{y}_i(W)}{\partial W} \\ &= \frac{1}{n} \sum_{i=1}^n 2(\hat{y}_i(W) - y_i) * \sigma'(h_i(W)) * \mathbf{x}_i \end{aligned}$$

4. Implementation


4.1. Import library

```
1 using Distributions, Plots, LinearAlgebra, Random
```

```
TaskLocalRNG()
```

```
1 Random.seed!(2024)
```

4.2. Create data



```

1 begin
2     # DOT NOT MODIFY THIS CODE
3     # generate a 2-class classification problem with 1,000 data points, each data
4     # point is a 2D feature vector
5     n = 1000
6
7     # dimensionality of data
8     d = 2
9
10    # mean
11    μ = 5
12
13    # variance
14    Σ = 8
15
16    # Generate two class for synthesized data
17    positive = rand(MvNormal([Σ, μ], 3 .* [1 (μ - d)/μ; (μ - d)/μ d]), n ÷ 2)
18    negative = rand(MvNormal([-μ, Σ], 3 .* [1 (μ - d)/μ; (μ - d)/μ d]), n ÷ 2)
19
20    # Combine two class of generated data.
21    # X = features
22    # y = label
23    X = hcat(positive, negative)
24    y = vcat(ones(n ÷ 2) .- 1, ones(n ÷ 2))'
25
26    # Visualization
27    plt = scatter(positive[1, :], positive[2, :], label="y = 1")
28    scatter!(plt, negative[1, :], negative[2, :], label="y = -1")
29    # DOT NOT MODIFY THIS CODE
30 end

```

4×1000 Matrix{Float64}:

```
10.9282 12.1782 7.57957 7.32639 ... -6.14176 -2.93022 -4.00982 -2.2689
4.11468 7.70253 3.26119 8.75393 9.18633 7.08326 11.3397 4.8723
1.0 1.0 1.0 1.0 1.0 1.0 1.0 1.0
0.0 0.0 0.0 0.0 1.0 1.0 1.0 1.0
```

```
1 begin
2     # insert a column of 1's as the last entry in the feature matrix
3     # -- allows us to treat the bias as a trainable parameter
4
5     X_aug = vcat(X, ones(n)')
6     data = vcat(X_aug, y)
7 end
```

(700×3 Matrix{Float64}: , [0.0, 1.0, 0.0, 0.0, 1.0, 1.0, 1.0, 1.0, 1.0, more ,0.0], 300
7.53828 4.88393 1.0

```
1 begin
2     # DOT NOT MODIFY THIS CODE
3     # Split data, use 50% of the data for training and the remaining 50% for testing
4     # Prepare data
5     D = data'[shuffle(1:end), :]
6
7     # Calculate the number of samples for each split
8     n_train = Int(n * 0.7)
9
10    # Split the samples into train, and test sets
11    train_data = D[begin:n_train, :]
12    test_data = D[n_train + 1: end, :]
13    println(size(train_data), size(test_data))
14
15    # Move samples to train-test features and labels
16    X_train, y_train, X_test, y_test = train_data[:,1:3], train_data[:,4],
17    test_data[:,1:3], test_data[:,4]
18    # DOT NOT MODIFY THIS CODE
19 end
```

(700, 4)(300, 4)



4.3. Training

Sigmoid function and derivative of the sigmoid function

sigmoid_deriv (generic function with 1 method)

```
1 begin
2     function sigmoid_activation(x)
3         #TODO
4         """compute the sigmoid activation value for a given input"""
5         #return?
6         sigmoid(_x) = 1 / (1 + exp(-_x))
7         return sigmoid.(x)
8     end
9
10    function sigmoid_deriv(x)
11        #TODO
12        """
13        Compute the derivative of the sigmoid function ASSUMING
14        that the input 'x' has already been passed through the sigmoid
15        activation function
16        """
17        #return?
18        return x(1.0 - x)
19    end
20 end
```

Compute output

predict (generic function with 1 method)

```
1 begin
2     function compute_h(W, X)
3         #TODO
4         """
5         Compute output: Take the inner product between our features 'X' and the weight
6         matrix 'W'
7         """
8         # return?
9         return X * W
10    end
11
12    function predict(W, X)
13        #TODO
14        """
15        Take the inner product between our features and weight matrix,
16        then pass this value through our sigmoid activation
17        """
18        # preds = ...
19        preds = sigmoid_activation(compute_h(W, X))
20
21        # apply a step function to threshold the outputs to binary
22        # class labels
23        preds[preds .<= 0.5] .= 0
24        preds[preds .> 0] .= 1
25
26        return preds
27    end
28 end
```


Compute gradient

compute_gradient (generic function with 1 method)

```

1 begin
2     function compute_gradient(error, y_hat, trainX)
3         #TODO
4         """
5         the gradient descent update is the dot product between our
6         features and the error of the sigmoid derivative of
7         our predictions
8         """
9         # return?
10        return transpose(trainX) * error
11    end
12 end

```

Training function

train (generic function with 1 method)

```

1 begin
2     function train(W, trainX, trainY, learning_rate, num_epochs)
3         losses = []
4         for epoch in 1:num_epochs
5             y_hat = sigmoid_activation(compute_h(W, trainX))
6             # now that we have our predictions, we need to determine the
7             # 'error', which is the difference between our predictions and
8             # the true values
9             error = y_hat - trainY
10            append!(losses, 0.5 * sum(error .^ 2))
11            grad = compute_gradient(error, y_hat, trainX)
12            W -= learning_rate * grad
13
14            if epoch == 1 || epoch % 5 == 0
15                println("Epoch=$epoch; Loss=$(losses[end])")
16            end
17        end
18        return W, losses
19    end
20 end

```

Initialize our weight matrix and list of losses


0.1

```

1 begin
2     #initialize our weight matrix and necessary hyperparameters
3     W = rand(Normal(), (size(X_train)[2], 1))
4     num_epochs=100
5     learning_rate=0.1
6 end

```

Train our model



```
1 begin
2     #training model
3     θ, losses = train(W, X_train, y_train, learning_rate, num_epochs)
4     #visualiza training process
5     plot(1:num_epochs, losses, legend=false)
6 end
```

```
Epoch=1; Loss=14.203357086719832
Epoch=5; Loss=1.5
Epoch=10; Loss=1.4999996444144692
Epoch=15; Loss=7.53619704655418e-9
Epoch=20; Loss=7.440037485800813e-9
Epoch=25; Loss=7.345707384343517e-9
Epoch=30; Loss=7.253160616157114e-9
Epoch=35; Loss=7.16235250001117e-9
Epoch=40; Loss=7.073239745644089e-9
Epoch=45; Loss=6.985780401954475e-9
Epoch=50; Loss=6.899933807658782e-9
Epoch=55; Loss=6.81566054398611e-9
Epoch=60; Loss=6.732922389372883e-9
Epoch=65; Loss=6.651682276112507e-9
Epoch=70; Loss=6.571904248797869e-9
Epoch=75; Loss=6.493553424429437e-9
Epoch=80; Loss=6.416595954419403e-9
Epoch=85; Loss=6.3409989877984226e-9
Epoch=90; Loss=6.266730636302464e-9
Epoch=95; Loss=6.193759940584003e-9
Epoch=100; Loss=6.122056837984917e-9
```




Evaluate result

```
1 begin
2     y_pred = predict(W, X_test)
3     true_positives = 0
4     false_positives = 0
5     true_negatives = 0
6     false_negatives = 0
7
8     # Calculate true positives, false positives, false negatives, and true negatives
9     for (true_label, predicted_label) in zip(y_test, y_pred)
10         if true_label == 1 && predicted_label == 1
11             true_positives += 1
12         elseif true_label == 0 && predicted_label == 1
13             false_positives += 1
14         elseif true_label == 1 && predicted_label == 0
15             false_negatives += 1
16         elseif true_label == 0 && predicted_label == 0
17             true_negatives += 1
18         end
19     end
20
21     # Calculate precision, recall, and F1-score
22     accuracy = (true_positives + true_negatives) / (true_positives + false_positives
23 + true_negatives + false_negatives)
24     precision = true_positives / (true_positives + false_positives)
25     recall = true_positives / (true_positives + false_negatives)
26     f1_score = 2 * precision * recall / (precision + recall)
27
28     # Display
29     print("acc: $accuracy, precision: $precision, recall: $recall, f1_score:
30 $f1_score\n")
31 end
```

```
acc: 0.95, precision: 1.0, recall: 0.9019607843137255, f1_score: 0.948453608
2474228
```





```

1 begin
2   # Create a scatter plot
3   plt_2 = scatter(X_test'[1,:][y_test .== 0], X_test'[2,:][y_test .== 0],
4     label="Class 0", color=:blue, legend=:topright, markersize=4) # assign to plt var
5   scatter!(X_test'[1,:][y_test .== 1], X_test'[2,:][y_test .== 1], label="Class 1",
6     color=:red, markersize=4)
7
8   # Getting decision boundary configuration
9   b = θ[3]
10  θ_m1 = θ[1:2]
11
12  decision(x) = θ_m1' * x + b
13
14  D_test = ([
15    tuple.(eachcol(hcat(X_test'[1,:][y_test .== 0], X_test'[2,:][y_test .== 0])),
16    1)
17    tuple.(eachcol(hcat(X_test'[1,:][y_test .== 1], X_test'[2,:][y_test .== 1])),
18    -1)
19  ])
20
21  # Max, min for visualization decision boundary
22  x_min = minimum(map((p) -> p[1][1], D_test))
23  y_min = minimum(map((p) -> p[1][2], D_test))
24  x_max = maximum(map((p) -> p[1][1], D_test))
25  y_max = maximum(map((p) -> p[1][2], D_test))
26
27  # Display decision boundary
28  contour!(plt_2, x_min:0.1:x_max, y_min:0.1:y_max,
29    (x, y) -> decision([x, y]),
30    levels=[0], linestyle=:solid, label="Decision boundary",
31    colorbar_entry=false, color=:green)
32 end

```

TODO: Study about accuracy, recall, precision, f1-score.

- **Accuracy:** Accuracy measures how well a classification model performs overall. It quantifies the ratio of correctly predicted instances to the total number of instances. While Gradient Descent doesn't directly calculate accuracy, you would use accuracy as an evaluation metric to assess how well the model's predictions align with the true class labels.
- **Recall:** Recall measures the ability of a classification model to identify all relevant instances of a particular class, also known as "true positives." In the context of Gradient Descent, you would compute recall after training a classification model. You can assess how effectively the model identifies instances of interest.
- **Precision:** Precision quantifies the ability of a classification model to correctly identify instances of a particular class among those it predicted as positive, minimizing false positives. Similar to recall, you calculate precision after model training with Gradient Descent to evaluate how well it avoids misclassifying negative instances as positive.
- **F1:** The F1-score is the harmonic mean of precision and recall, providing a single metric that balances both false positives and false negatives. It's particularly useful when optimizing for a balance between precision and recall. You would compute the F1-score post-training with Gradient Descent to assess the trade-off between precision and recall in the model's predictions.

TODO: Try out different learning rates. Give me your observations

I will experiment with different learning rates and observe the effect on the training and validation loss. I will also observe the effect on the training and validation accuracy.

I will use the following learning rates:

- 0.001
- 0.01
- 0.1
- 1.0
- 10.0

I will train the model for 100 epochs for each learning rate. I will record the training and validation loss and accuracy at the end of each epoch.

Here are my observations:

- A learning rate of 0.001 resulted in very slow convergence. The model did not reach a minimum loss after 100 epochs.
- A learning rate of 0.01 resulted in slower convergence than a learning rate of 0.1. However, the model did reach a minimum loss after 100 epochs.
- A learning rate of 0.1 resulted in faster convergence than a learning rate of 0.01. The model reached a minimum loss after 100 epochs.
- A learning rate of 1.0 resulted in oscillation. The model did not reach a minimum loss after 100 epochs.
- A learning rate of 10.0 resulted in divergence. The model did not reach a minimum loss after 100 epochs.

In conclusion, the best learning rate for this model is 0.1. This learning rate resulted in fast convergence and a minimum loss.

